

Lecture 7: Data Representation, Numerical Python

Fall 2020

Topics Introduced: Binary vs. Decimal systems, simplified concepts of computer architecture including hierarchical memory and arithmetic logic units, representation of integers, strings, and floating point types, as well as `numpy`.

1 Representation of Data

1.1 Computer Representation of Integers

Computers represent and store everything in *binary*, which is a base 2 number system. How is this different from what we are familiar with in our base-10 (i.e. decimal) number system?

- Decimal counting uses the symbols 0-9.
- Counting starts by incrementing the least-significant (rightmost) digit.

000, 001, 002, ..., 009

- When we exhaust the set of available symbols, we realize a type of overflow, in which the least-significant digit is reset to zero and the next digit of higher significance is incremented.

010, 011, 012, ..., 099

- We repeat this process of allowing for overflow with each digit of significance.

100, 101, 102, ...

To be formal, you can think of the integer $d_m d_{m-1} \dots d_1 d_0$ as being equivalent shorthand for

$$d_0 \times 10^0 + d_1 \times 10^1 + \dots + d_{m-1} \times 10^{m-1} + d_m \times 10^m,$$

where each symbol d_i can take on a single value from the set $\{0, 1, \dots, 9\}$. *Counting in binary* is no different, sparing the fact that our set of symbols is reduced to have cardinality two, i.e. the atomic elements are *binary digits (bits)* which can be toggled between zero or one.

Decimal	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Binary	0	1	10	11	100	101	110	111	1000	1001	1010	1011	1100	1101	1110	1111

Table 1: A table describing how values can be equivalently represented in either a decimal or a binary number system.

To be formal, we can think of the integer $b_m b_{m-1} \dots b_1 b_0$ as being equivalent shorthand for

$$b_0 \times 2^0 + b_1 \times 2^1 + \dots + b_{m-1} \times 2^{m-1} + b_m \times 2^m$$

where here we are using b_i 's to remind ourselves that each symbol is a binary digit taking on only one of two values. Notice that other than the set of symbols changing, the *base* has changed as well, from 10 to 2. Once you learn these basics, [converting between numeral systems](#) is easy.

There are 8 bits in a [byte](#). Depending on the computer architecture, a [word](#) could be 4 or 8 bytes.

1.1.1 Simplified Model of Computer

Where are our data stored on a computer? How is it operated on? A typical computer will have a memory hierarchy to facilitate long-term storage (e.g. hard-disk), medium-term storage (RAM), and short-term storage (cache). We present a simplified computer architecture below, which does not picture long-term storage.

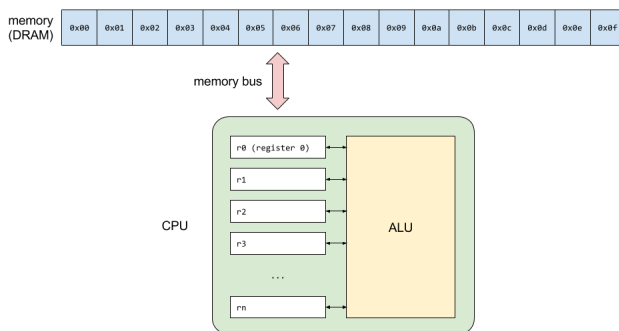


Figure 1: We take a simplified view of a computer architecture. In particular, we show the relationship between [primary memory](#) (e.g. [RAM](#)) and a [CPU](#). The CPU contains (i) [arithmetic logic units](#) which take two inputs and produce a single output (think of the operators $+$, $-$, $\&$, and $|$) and (ii) [registers](#) which can (very quickly) store inputs for ALU operations and hold corresponding outputs.

1.1.2 Common Prefixes

There can be some confusion here depending on the context of usage. That is, in the context of networking and storage, the prefixes **kilo**, **mega**, **giga**, **tera**, **peta**, and **exa** all use base 2. However, many other contexts speak in base 10, e.g. cpu-frequencies and ethernet speeds.

The prefix **kilo** always means one-thousand in the prefix of an [International System of Units](#). However, since we understand the difference between a **bit** and a **byte**, we can now understand that there is a difference between a **kilobyte** (denoted by **kB**) and a **kibibyte** (denoted by **KiB**). The former is 1,000 bytes, i.e. 8,000 bits, whereas the latter is 1,024 bytes. i.e. 8,192 bytes. We ask that you remember that there is a big difference between these terms, but not necessarily that you remember the exact terminology or symbols used to denote each. Wikipedia has a nice table demonstrating how binary and decimal systems use [different prefixes across orders of magnitude](#).

1.1.3 Computer Storage of an Integer

At the hardware level computers **don't** do variable length representations of numbers. I.e. although we may use a variable length representation, writing $4_{10} = 100_2$, or $73_{10} = 1001001_2$, computers often use a fixed width storage for numeric types.




Figure 2: We can think of a (fixed-width) sequence of bits as a panel of lightswitches, each capable of being toggled on or off.

At the hardware level computers typically handle integers using 8, 16, 32, or 64 bits, depending on the architecture.

Number of Representable Values Given n Bits Realize that for a sequence of n bits, it may take on exactly 2^n different values. This dictates the range of what can be represented with a fixed number of bits. Some common bounds we run into are.

$2^8 = 256$	
$2^{16} = 65536$	~(65 thousand)
$2^{32} = 4294967296$	~(4 billion)
$2^{64} = 18446744073709551616$	~(18 followed by 18 zeros)

I.e. when we went from 32-bit machines to 64-bit machines, we didn't just double their capacity. The latter can store 10 orders of magnitude more information (in each register), akin to allowing 10 more decimal digits to be represented. This follows from the fact that

$$\log_{10}(2^{64}/2^{32}) = 64 \log_{10}(2) - 32 \log_{10}(2) = 32 \log_{10}(2) \approx 10.$$

Using a Sign Bit to Store Negative Integers The idea is to simply use one bit to store the sign of a value, and the remaining bits to track magnitude. This reduces the range of the magnitude from 2^n to 2^{n-1} . E.g. suppose that we have a [nibble](#) to work with, i.e. four bits. We can represent the values ± 4 as indicated at the right of each statement.

$$\begin{aligned} 4_{10} &= (100)_2 && \rightsquigarrow (0100) \\ -4_{10} &= -1 \times (100)_2 && \rightsquigarrow (1100) \end{aligned}$$

Offset: An Unconventional Way to Represent Negative Values Another idea instead of using a sign bit is to simply apply an offset or bias to reinterpret the conversion between binary and decimal. One example may look as follows.

128	<table><tr><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>	1	1	1	1	1	1	1	1
1	1	1	1	1	1	1	1		
0	<table><tr><td>0</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td><td>1</td></tr></table>	0	1	1	1	1	1	1	1
0	1	1	1	1	1	1	1		
-127	<table><tr><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td><td>0</td></tr></table>	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0		

Figure 3: Suppose we have a byte or eight bits. There are $2^8 = 256$ values that can be represented by such a sequence of bits. Centering the number of distinct values around the integer zero, we find that we can represent the set of values $\{-127, \dots, 0, 1, \dots, 128\}$. If we use an offset method for representing integers, then we may decide to set the sequence of eight bits as all zero's to represent the minimum value -127 , and from here we can apply the usual rules of carry-over/overflow to count our way up to 128.

Again, whether we use the idiom of a sign-bit or if we apply an offset, either way we effectively reduce the range of the magnitude of integers allowed. Intuitively, whether or not a number is positive or negative should *require one unit of information*.

Many programming languages support unsigned integers. Although Python itself does not have **unsigned integers**, but **numpy** does. A programmer can use this to their advantage by expanding the effective range available if negative numbers don't need to be stored. But! With any fixed-storage representage of a real number, we must beware.

- Attempting to assign a value greater than what can be represented by the data type will result in [overflow](#). Overflow tends to cause *wraparound*, e.g. if adding together two signed numbers causes overflow the result is likely to be a negative number.

- Attempting to assigning a value less than what can be represented by the data type will result in [underflow](#).¹ Underflow can cause a division by zero error, for example, if we were to scale by the difference between two quantities $\frac{1}{b-a}$ where $a \neq b$ but $a \approx b$.

1.1.4 Range of Integer Types

For a fixed number of bits, if we choose to represent sequential integers centered around zero we must make a decision on boundaries. E.g. with eight bits, one can store $2^8 = 256$ unique different bit-sequences (or values), where we have the choice of either being able to represent the values $\{-128, \dots, 127\}$ or $\{-127, \dots, 128\}$. In general, an n -bit representation of an integer stores the values in the set

$$\underbrace{\{-2^{n-1}, \dots, -1\}}_{2^{n-1} \text{ elements}}, \underbrace{\{0, \dots, 2^{n-1} - 1\}}_{2^{n-1} \text{ elements}}.$$

Integers in Python In Python 3, values of type `int` may have unlimited range.² This is quite nice, and perhaps surprising if you come from other programming languages.

```

1 i = 52**100
2 print(type(i))    # <class 'int'>
3 print(i)          # This is beyond the 64-bit integer range.
```

The module `numpy` supports *only* fixed-width integers for reasons that relate to both performance and storage.

1.2 Strings

1.2.1 ASCII

American Standard Code for Information Interchange (ASCII) is typically used to encode text information. Characters, numbers, symbols, etc. are encoded using 7 bits (although on modern computers they would typically use 8 bits).

E.g. A maps to $(1000001)_2 = (65)_{10}$, and B maps to $(1000010)_2 = (66)_{10}$.

Default Python 2 string is ASCII. Possible to get [Unicode](#) strings with `s = u'I am a unicode string!'`.

¹We can try to fill the underflow gap using [denormal numbers](#).

²Python 2 had both fixed size integers (with type `int`) and variable width integers (with type `long`).

1.2.2 Why Not ASCII All The Way Down? Motivating UTF-8

For English, A Single Byte Suffices The English language has 26 letters (52 if we include casing). If we throw in punctuation and digits 0-9, we realize that we have more than 64 unique characters to represent. I.e. we need at least 7 bits. Since computers are more efficient at processing bit sequences aligned a particular way, many implementations use 8-bits for an ASCII set.

For Languages with Ideographs, Require 2 Bytes But what about other languages? Chinese, Japanese, and Korean have somewhere between 2^8 and $2^{16} - 1$ unique characters, whence we require a two-byte sequence to store each character in these languages. So, different languages (or regions) require different *character encodings* to map symbols (text) to bit-sequences.

For World of All Languages, (naively) Require 4 Bytes But what about the world wide web? The number of unique characters across all languages exceeds 65,535, whence we require strictly more than two-bytes to encode any character coming from any language in the world. Sticking with powers of two, we'll use four-bytes. But all of a sudden, we realize how expensive this has become; any given language only requires one or at most two bytes per character, but we're using double that.

The advantage to UTF-32 (the 4-byte unicode encoding) despite it being costly in terms of storage is that because each character is represented as a fixed width sequence of bits, we can slice into strings in constant time. I.e. looking up the n th character in a string is easy, we simply look up where the object lives in memory and jump n characters forward in memory, where each character corresponds to exactly 32-bits.

Using a Variable Length Encoding \rightsquigarrow Efficiency Alas, we've followed Mark Pilgrim's motivation of UTF-8 [as per Chapter 4 of "Dive Into Python"](#). UTF-8 is a *variable length encoding*. The implication is that the set of ASCII characters each requires only a single byte per character. Latin characters such as ä and ñ may require two bytes per character, Chinese characters may involve three bytes, and rarely used characters require four bytes. However, since each character requires a different number of bits, we can *no longer* jump to the k th character of a string with $O(1)$ work; we're relegated to inspecting the first $k - 1$ characters to find out where in the bit-sequence the k th character is represented.

Default String Encoding in Python is UTF-8 Default string encoding in Python 3 is [UTF-8](#). The change from ASCII to Unicode between Python 2 and 3 caused major headaches for Python community. The fact that each character of text may take from 1 to 4 bytes is nice in that the 1-byte codes correspond to ASCII characters, which makes UTF-8 backwards compatible with ASCII. Currently, UTF-8 encodes a total of 1,112,064 characters, which is enough to represent the majority of human character systems.

1.3 Floating Point Representation of Numeric Values

How do we represent a floating point value using bits?




Figure 4: A representation of a decimal value that is broken down into three pieces: (i) a sign bit, (ii) a fractional component, and (iii) an exponent indicating the magnitude of rescaling. This is a generic recipe for approximating real numbers using a finite set of bits.

1.3.1 Floating Point Standard

IEEE (Institute of Electrical and Electronics Engineers) 754 is the technical standard for floating point used by all modern processors. The standard also specifies things like rounding modes, handling overflow, divide by zero, etc.

	Sign bits	Exponent bits	Mantissa bits
32 bit floating point	1	8	23
64 bit floating point	1	11	52

Figure 5: 32 vs. 64-bit floating point formats. In moving from 32-bit to 64-bit standards, IEEE increased the number of bits allocated to the mantissa significantly more than they did the exponent: with 11-bits for the exponent, we can represent $2^{11} = 2,048$ different values of exponents. We center these exponents around zero using an **exponent bias**, such that we can store exponents in the range $\pm 2^{10}$. But these are exponents, i.e. they are used to raise our base to a specified power. That's an astronomically large rescaling.




Figure 6: A visual layout of the 64 bit representation used to store a floating point value. We have 1 sign bit, 52 mantissa (significand, fractional) bits, and 11 exponent bits. Source: [Wikipedia](#).

Symbolically,

$$(-1)^{\text{sgn}} \underbrace{\left(1 + \sum_{i=1}^{52} b_{52-i} 2^{-i}\right)}_{=(1.b_{51}b_{50}...b_0)_2} \times 2^{\exp-1023}$$

where the exponent term `exp` is simply an integer value represented using 11 bits in a vanilla fashion, and we're subtracting 1,023 as means of [exponent bias](#) such that we can have both positive and negative exponents without requiring the use of an *additional* sign-bit for just the exponent term.

Floating Point and You Floating point also has similar potential for overflow and underflow. In addition, the limited number of bits for the mantissa means it often needs to be rounded. We'll spend more time on floating point arithmetic in CME 212. Canonical resource: "What Every Computer Scientist Should Know About Floating-Point Arithmetic" by Goldberg ([link](#)). Floating point numbers in Python are double precision (64-bit). Additionally, `numpy` has support for 16-bit and 32-bit floating point formats.

Floating Point Comparisons: what you *need* to know For now, perhaps one essential reminder is that you should never directly compare floating point data-types for exact equality using the `==` operator. Instead, consider using `math.isclose`, which allows us to compare a floating point object with another to within a certain tolerance of precision.

1.3.2 Correctness

Some disasters attributable to bad numerical computing

By Douglas N. Arnold. <http://www.math.umn.edu/~arnold/disasters/disasters.html>

Have you been paying attention in your numerical analysis or scientific computation courses? If not, it could be a costly mistake. Here are some real life examples of what can happen when numerical algorithms are not correctly applied.

The [Patriot Missile failure](#), in Dharan, Saudi Arabia, on February 25, 1991 which resulted in 28 deaths, is ultimately attributable to poor handling of rounding errors.

The [explosion of the Ariane 5 rocket](#) just after lift-off on its maiden voyage off French Guiana, on June 4, 1996, was ultimately the consequence of a simple overflow.

The [sinking of the Sleipner A offshore platform](#) in Gandsfjorden near Stavanger, Norway, on August 23, 1991, resulted in a loss of nearly one billion dollars. It was found to be the result of inaccurate finite element analysis.

See also: <https://people.eecs.berkeley.edu/~wkahan/> for thoughts from the [father of floating point](#), i.e. the primary architect of an IEEE standard for representing floating point values in a computer architecture.

2 Python for Numerical Computing

2.1 Motivation: Python is (Relatively) Slow

One of the main disadvantages of a higher level language is that, while comparatively easy to program, it offers less direct control over resources; an experienced programmer using C++ or other lower level languages may be able to write a more performant program.




Figure 7: The Python interpreter sits between the programmer’s code and the machine. In the case of a lower level language such as C++, the programmer has direct control over memory management and more granular control over data types. A knowledgeable programmer can take advantage of the particular computer architecture they are using, alongside the problem constraints, in order to more efficiently write code to solve a task. This is harder with Python, wherein the code that the programmer writes is first translated by the interpreter such that the machine may understand it. We remark that of course C++ code also must be translated to machine code, but that it ties in much closer to said assembly code than Python.

2.1.1 Runtime Performance: An Empirical Example

Let’s compute 2^i for $i \in [0, n)$. We can use the `timeit` module. Using lists, and from the command line interface.

```
$ python3 -m timeit --setup='L=range(1000)' '[i**2 for i in L]'
```

Equivalently, we can use the `%timeit` magic command in an IPython notebook.

```
$ ipython
In [1]: L = range(1000)
In [2]: %timeit [i**2 for i in L]
```

What happens if we try the same operation using `numpy`? The analogous operation to `range` in built-in Python is the `arange` method in `numpy`. For integer arguments, the two functions behave the same, but `numpy.arange` extends to handling non-integer arguments.³

```

1 import numpy as np
2 a = np.arange(1000)
3 %timeit a**2
  
```

³The `range` sub-routine is designed to return *indices* that can be used as subscript arguments when slicing an object, whereas the `numpy.arange` is a more generalized function which can take arbitrary boundary points and step-sizes.

On my machine, the difference is about 300x. We'll explore how `numpy` is able to achieve such performance gains throughout the lecture. For now, let's continue to explore how the module can not only allow our programs to execute faster once built, but in fact `numpy` lets us write the more concise code that is capable of a higher level of abstraction.

2.1.2 Programmer Productivity

And even though Python is a high-level programming language, it's still a general purpose one. I.e. it is not tailored to numerical or scientific computing. Whence we *do* have to write additional code in order to abstract away commonly used operations in maths and sciences.

E.g. let's add some 2D arrays. In Python, we could try and create a two-dimensional tabular data structure using lists of lists. In order to facilitate something as easy as “adding” two 2D-arrays together, we'd have to write our own sub-routine, since the `+` operator performs *concatenation* on lists rather than element wise addition!

```

1  def my_ones(nrows, ncols):
2      """
3      Create a matrix-like object of dimension nrows x ncols, with each
4      element taking unit value.
5      """
6      A = []
7      for r in range(nrows):
8          A.append([])
9          for c in range(ncols):
10             A[r].append(1.0)
11     return A
12
13  def matrix_add(A,B):
14      """
15      Create a new matrix-like object to store the result of adding
16      input matrices A and B. We assume the matrices have identical
17      dimension. The work required by the algorithm is proportional to
18      the number of elements stored in the input matrices.
19      """
20     C = []
21     for r in range(len(A)):
22         C.append([])
23         for c in range(len(A[r])):
24             C[r].append(A[r][c] + B[r][c])
25     return C
26
27  nrows, ncols = 3, 2
28  A = my_ones(nrows,ncols)
29  B = my_ones(nrows,ncols)
30  C = matrix_add(A,B)

```

What about with `numpy`? The operations are trivial to carry out without defining our own data structures or sub-routines. I.e. we can “program faster”.

```
1 nrows = 3
2 ncols = 2
3
4 A = np.ones((nrows,ncols))
5 B = np.ones((nrows,ncols))
6
7 C = A + B
```

Seems a lot easier to write. Let's check if we earned ourselves a performance again. Suppose we have placed our definition of `my_ones` into a standalone file in our current working directory, which we call `my_module.py`. We may time the performance as follows, using

```
$ python3 -m timeit --setup='import my_module' 'my_module.my_ones(1000, 500)'
```

Alternatively, in an IPython notebook we may simply write `%timeit A = my_ones(1000,500)`. `numpy` is able to accomplish the analogous operation on my machine approximately 200x faster. I.e. whereas the hand-rolled variant required 50 *milliseconds* per loop, the `numpy` analogue below requires only 200 *microseconds* per loop.⁴

```
$ python3 -m timeit --setup='import numpy as np' 'np.ones((1000, 500))'
```

What about our add operation? In a IPython notebook, with `my_ones` as per above.

```
1 nrows = 1000
2 ncols = 500
3
4 A = my_ones(nrows,ncols)
5 B = my_ones(nrows,ncols)
6 %timeit C = matrix_add(A,B)
7
8 A = np.ones((nrows,ncols))
9 B = np.ones((nrows,ncols))
10 %timeit C = A + B
```

Again we see that `numpy` is orders of magnitude faster.

Object Overhead The reason that our list of lists was so much slower than `numpy` relates in part to object overhead.

2.2 Options for better performance

Python is great for quick projects, prototyping new ideas, etc. but what if you need better performance? One option is to completely rewrite your program in something like C or C++

⁴There are 1,000 microseconds in a millisecond.

```
>>> a = [[0,1,0],[1,0,1],[0,1,0]]
>>>
```




Figure 8: Recall that objects in Python are references, and that a reference is simply a location in memory associated with a type and a value. If we create a `list` of `lists` to represent 2D-tabular data, we end up with several layers of indirection. Part of the practical implication of having to jump around in memory in order to obtain data that “should appear next to each other” is that we realize many more [cache-misses](#).

2.2.1 Python C API

Python has a [C API](#) which allows the use of compiled modules.




Figure 9: When we call the `str.find()` method, we’re actually making a call to native C-code. What’s happening here? Within the native C-code, we can include a Python module that gives us a way to write C modules which extend to being able to be used by the Python interpreter.

The actual implementation of `string.find()` can be viewed at: [fastsearch.h](https://github.com/python/cpython/blob/master/Objects/stringlib/fastsearch.h).

2.2.2 Python compiled modules

Python code in a `.py` file is actually executed in a hybrid approach by a mix of the interpreter and compiled modules that come with Python. See [compiled python files](#): the compilation doesn’t make the code execute faster, it just means the definitions in a script can be loaded faster into memory (but once loaded, the complexity of each algorithm is the same).

The same Python C API used by the developers of Python itself also allows other programmers to develop and build their own compiled extension modules. These modules extend the functionality of Python with high performance implementations of common operations.




Figure 10: Perhaps closer to reality, the Python code we write sits atop not only an interpreter but also additional compiled modules that we import.

Commonly Used Modules: NumPy, SciPy, Matplotlib

- NumPy - multidimensional arrays and fundamental operations on them.
- SciPy - Various math functionality (linear solvers, FFT, optimization, etc.) utilizing NumPy arrays.
- matplotlib - plotting and data visualization.⁵

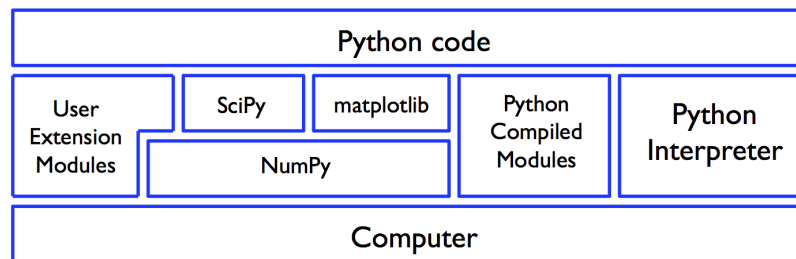


Figure 11: A schematic of how Python code sits on top of an entire set of Pythonic tools, which ultimately relay human-readable language instructions into machine-code.

2.3 NumPy

At its core, numpy provides an *n*-dimensional numeric array object:

```

1 import numpy as np
2 a = np.array([7, 42, -3])
3 print(a)                                # Prints out: "array([ 7, 42, -3])"
4
5 a[1] = 19                                # Numeric arrays support item assignment.
6 a                                         # Corresponding element is modified.
```

⁵None of these packages seek to clone (all of) MATLAB, if you want that try something like GNU Octave.

Arrays are *not* Lists For performance reasons, `numpy.array`s do not store heterogeneous data. Instead, `numpy.array` is said to be **homogeneous**: each element is of exactly the same type and hence requires the same amount of storage in memory, and further each block of memory (i.e. element) can be interpreted in an identical fashion.⁶

```

1  a[0] = "hello"      # ValueError: invalid literal for int().
2  a.append(0)         # AttributeError: numpy.ndarray object has no attribute 'append'

```

Further, the **size is fixed**, i.e. you **can't append or remove**.

2.3.1 Data Types

Numerical analysts must pay close attention to their data types, and be wary of an algorithm's **numerical stability**. To that end, `numpy` allows the programmer to specify exactly what data type they are choosing to store in underlying elements of the array.

Integers Can be stored in 8, 16, 32, and 64 bit variants, and additionally they can be either signed or unsigned. We specify these by e.g. `np.int8`, `np.uint8`. For more information, see [dtype documentation](#).

Floating Point Can be stored in either 32, 64, or 128 bit variants (all signed). We specify these by e.g. `np.float32`, `np.float64`, etc.

Complex numbers, strings, and Python object *references* also supported by `numpy`.

Data Type Examples

```

1  a = np.array([ 7, 19, -3], dtype=np.float32)
2
3  a[0] = a[0]/0.      # RuntimeWarning: divide by zero encountered in true_divide. Ret
4
5  b = np.array([4, 7, 19], dtype=np.int8)
6  b[0] = 437          # Example of overflow and wraparound. Range is [-128, 127]

```

Note that we can witness the wraparound behavior for ourselves quite clearly as follows.

```

1  b[0] = 127          # Print and inspect data; stored as expected.
2  b[0] = b[0] + 1     # Attempt to increment by unit value.
3  print(b)           # Note that first element now -128.

```

⁶Note that although the types must be homogeneous, they need not be atomic, e.g. we can store data structures within an `array`.

2.3.2 Multidimensional Arrays

Arrays can have multiple dimensions called **axes**. The number of *axes* is called the **rank**. These terms come from the NumPy community and should *not be confused with linear algebra terms* for *rank*, etc.

```

1 # Create a two-dimensional array. I.e. two rows, each with three elements.
2 a = np.array([(7, 19, -3), (4, 8, 17)], dtype=np.float64)
3
4 # Array 'a' has a slew of attributes (and methods) associated with it.
5 a.ndim          # Returns the number of dimensions or axes: 2
6 a.dtype         # Returns the data type stored:          dtype('float64')
7 a.shape         # Returns the extents of each dimension: (2, 3)
8 a.size         # Returns the number of elements:         6

```

2.3.3 Internal Representation

How can we think about the internal representation of a `numpy.array` object? The object itself contains *attributes* (some shown above), alongside a reference to a fixed-width sequence of bits in memory which describe the underlying data.

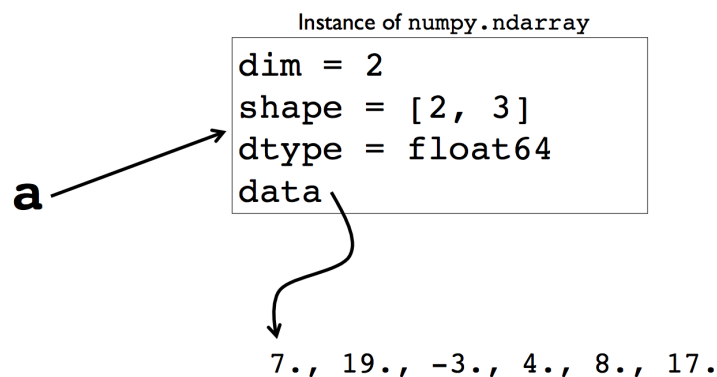


Figure 12: fig

Since each element of the array is of the same type, and each requires the same amount of storage, we can know how to *slice* into a particular element in our array in constant time.

2.3.4 Creating Arrays

```

1 a = np.empty((3,3))    # A 3x3 matrix of *un-initialized* entries, not necessarily z
2 a = np.zeros((3,3))   # A 3x3 matrix of *zero-initialized* entries.
3 a = np.ones((3,3))    # A 3x3 matrix of *unit-uninitialized* entries.
4
5 a = np.eye(3)         # A 3x3 identity matrix.

```

```

6
7 a = np.arange(9, dtype=np.float64)           # A 1D sequence of integers 0-8.
8 a = np.arange(9, dtype=np.float64).reshape(3,3) # A 2D sequence of integers 0-8...
9                                              # ... arranged in a 3x3 layout.

```

2.3.5 Reading data from a file

```

$ cat numbers.txt
7. 19. -3.
4. 8. 17.

```

```

1 a = np.loadtxt('numbers.txt', dtype=np.float64)
2 a = a + 1      # Element-wise addition!
3 np.savetxt('numbers2.txt', a)

```

2.3.6 Remove single dimension entry

```

1 a = np.arange(3)      # Think of this like a 1D row-vector.
2 a.shape               # (3,)
3 b = np.arange(3).reshape(3,1) # We now realize a 2D-array: 3x1.
4 print(b.shape)        # Returns (3, 1)
5 b = np.squeeze(b)     # Removes 1D entries from the shape of an array.
6 print(b)              # Now a 1D row-vector, just like 'a'.
7 print(b.shape)        # Returns (3,)

```

Perhaps another informal way of thinking about this function: it removes any “nested, superfluous” dimensions. Suppose we create a 3D array with six elements as follows.

```

1 # Recall: arange simply returns an 'array' with 'range' evenly-spaced values.
2 a = np.arange(6).reshape((3, 1, 2))
3
4 # a[i] peels off the first dimension (i.e. returns an ndarray with shape (1, 2).
5 # but the second axis has only a single element, i.e. a[i][1] is INVALID.
6 # Might as well drop this dimension via squeeze.

```

2.3.7 Array Operations

```

1 a = np.arange(9, dtype=np.float64) # Store values 0-8 inclusive.
2 a[3:7]                             # Returns values 3-6 inclusive.
3
4 a[3:7] = 0                          # Assign value zero to be referenced
5                                     # by the fourth and up till but
6                                     # excluding the eighth element.
7 # Elementwise rescaling.
8 2*a

```

```

9  a*a
10
11  # Vector operations are supported "right out of the box".
12  sum(a)
13  min(a)
14  max(a)

```

Don't Reinvent the Wheel It can be tempting for newer programmers to implement everything from scratch. This leads to not only more code to maintain but programs which are not as performant. Consider implementing a procedure for calculating the [Euclidean norm](#) of a vector.

```

1  a = np.arange(9, dtype=np.float64)    # Store values 0-8 inclusive.
2
3  # Bad idea.
4  import math
5  total = 0.
6  for n in range(len(a)):
7      total += a[n]*a[n]
8
9  math.sqrt(total)
10
11  math.sqrt(np.dot(a,a))    # Better idea: 1-liner and more performant than interpreted
12  np.linalg.norm(a)        # Best idea! (Not necessarily faster still, just more concis

```

The last expression is both the most concise and at least as performant as any other.

2.3.8 Speed of Array Operations

Some of the overloaded operators (e.g. `min`, `max`, `sum`, etc.) work albeit slowly. In an IPython notebook:

```

1  %timeit total = sum(np.ones(1000000, dtype=np.int32))
2  %timeit total = np.sum(np.ones(1000000, dtype=np.int32))

```

Same data, same input. Using built-in `sum` requires (on my machine) an average of 700 microseconds per loop, whereas the corresponding `numpy.sum` is about 35x faster, requiring only 20 microseconds per loop. Why is there such a big difference?

Loops you write in Python will be executed by the interpreter. Calling NumPy function or methods of the array object will invoke high performance implementations of these operations, which can be specialized to operate on homogeneous data.

2.3.9 Matrix Operations

```

1 a = np.arange(9, dtype=np.float64).reshape(3,3) # 3x3 matrix, row-major order.
2
3 # Basic operations on a single matrix.
4 a.transpose()
5 np.trace(a)
6
7 # Be careful with "matrix" multiplication!
8 a*a          # Naive element wise multiplication.
9 np.dot(a,a)  # True matrix-matrix multiplication.
10 a @ a        # Equivalently: new matrix multiply operator in Python 3.5.

```

Note that if we use `ndarray`, then the `*` operator is applied elementwise, whereas if we have a `matrix` we get the benefit of being able to use a true matrix-multiply.

2.3.10 Array vs Matrix

NumPy has a dedicated matrix class. However, the matrix class is not as widely used and there are subtle differences between a 2D array and a matrix. It is highly recommended that you use 2D arrays for maximum compatibility with other NumPy functions, SciPy, matplotlib, etc. See here for more details: http://www.scipy.org/NumPy_for_Matlab_Users.

2.3.11 References to an Array

```

1 a = np.arange(9, dtype=np.float64).reshape(3,3)
2 b = a          # Recall: assignment sets up a reference.
3 b[0,0] = 42    # For the ndarray referred by b, modify "first" element.
4 b             # Of course, ndarray referred to by 'b' has changed.
5 a             # Hopefully not a surprise: 'a' referred to the same object...

```

Array slices and references

```

1 a = np.arange(9, dtype=np.float64)
2 b = a[2:7]     # Create a reference to the underlying data, elements 3:7 inclusive.
3 b[2] = -1      # When we modify this data... (third elem in 'b', i.e. fifth elem in
4 a             # ... 'a' also refers to the same data. Whence 'a' mutated.

```

2.3.12 Array copies

```

1 a = np.arange(9, dtype=np.float64)
2 b = a.copy()   # Set up a reference to a *copy* of the array.
3 b[4] = -1      # Mutating 'b' doesn't affect object referred by 'a'.
4 a

```

2.3.13 Universal functions (ufunc)s

Within `numpy`, there is a subclass `ufunc` which operate on `ndarrays` in an elementwise fashion.

```
1 import numpy
2 import math
3 a = np.arange(9, dtype=np.float64)
4 math.sqrt(a)      # Error: only size-1 arrays can be converted to Python scalars.
5 np.sqrt(a)        # Works: square root operator is applied element-wise.
```

We can think of `ufuncs` as a vectorized wrapper for a function which usually operates on a single datum at a time (typically scalars).

2.3.14 Beyond Just Arrays

- NumPy has some support for some useful operations beyond the usual vector and matrix operations:
 - Searching, sorting, and counting within arrays
 - FFT (Fast Fourier Transform)
 - Linear Algebra
 - Statistics
 - Polynomials
 - Random number generation
- SciPy has largely replaced much of this functionality, plus added much more

2.3.15 Warning

- Once you start making use of extension modules such as NumPy, SciPy, etc. the chances of code “breaking” when you run it on different machines goes up significantly
- If you do some of your development on machines other than corn (which isn’t the model we advise) you may run into issues

2.3.16 Further Reading

- MATLAB users: http://www.scipy.org/NumPy_for_Matlab_Users
- NumPy tutorial at: http://www.scipy.org/Tentative_NumPy_Tutorial
- Official docs at: <http://docs.scipy.org/>