

Lecture 13 part 2

Fall 2020

Container iteration

Container iteration example 1

We can use C++11 for-loop style syntax.

src/iter1.cpp:

```
#include <iostream>
#include <vector>

int main()
{
    std::vector<double> vec;

    vec.push_back(7);
    vec.push_back(11);
    vec.push_back(42);

    // Creates a copy v for each element in vec and increments the copy
    for (auto v : vec)
        ++v;

    // The original elements of the vector vec are unchanged
    for (auto v : vec)
        std::cout << v << std::endl;

    return 0;
}
```

Output:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/iter1.cpp -o src/iter1
$ ./src/iter1
7
11
42
```

Container iteration example 2

We can also iterate over references to the underlying data.

src/iter2.cpp:

```
#include <iostream>
#include <vector>
```

```

int main()
{
    std::vector<double> vec;

    vec.push_back(7);
    vec.push_back(11);
    vec.push_back(42);

    // Creates a reference v to each element in vec and increments each element.
    for (auto& v : vec)
        ++v;

    // The original elements of the vector vec are incremented by one.
    // Here using constant reference to read vector elements.
    for (const auto& v : vec)
        std::cout << v << std::endl;

    return 0;
}

```

Output:

```

$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/iter2.cpp -o src/iter2
$ ./src/iter2
8
12
43

```

Map

- A C++ map is analogous to a dictionary in Python.
- However, since C++ is statically typed we need to specify data type for both the key and the value when instance is declared!

Our first map

src/map1.cpp:

```

#include <iostream>
#include <map>

int main()
{
    std::map<char, std::string> dir;

    dir['A'] = std::string("south");
    dir['B'] = std::string("north");
    dir['C'] = std::string("east");
    dir['D'] = std::string("west");

    std::cout << "dir[C] = " << dir['C'] << std::endl;
    std::cout << "dir[A] = " << dir['A'] << std::endl;

    return 0;
}

```

Output:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map1.cpp -o src/map1
$ ./src/map1
dir[C] = east
dir[A] = south
```

Map iteration

Iterating over a `map` using C++11 syntax results in obtaining a copy of a key-value pair each iteration. In the second for-loop below, we show how to obtain a *reference* to the key-value pair.

src/map2.cpp:

```
#include <iostream>
#include <map>

int main()
{
    // Define a map 'dir' with characters as keys and strings as values
    std::map<char, std::string> dir;

    dir['A'] = std::string("south");
    dir['B'] = std::string("north");
    dir['C'] = std::string("east");
    dir['D'] = std::string("west");

    // Printing by value (usually not a good idea)
    for (auto d : dir)
    {
        std::cout << "d[" << d.first << "] = " << d.second << std::endl;
    }
    std::cout << std::endl;

    // Printing by constant reference
    for (const auto& d : dir)
    {
        std::cout << "d[" << d.first << "] = " << d.second << std::endl;
    }

    return 0;
}
```

Output:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map2.cpp -o src/map2
$ ./src/map2
d[A] = south
d[B] = north
d[C] = east
d[D] = west

d[A] = south
d[B] = north
d[C] = east
d[D] = west
```

Older style iteration

It's also possible to use *iterators*, although this is more verbose and more old-school in terms of iteration paradigms.

src/map3.cpp:

```
#include <iostream>
#include <map>

int main()
{
    std::map<char, std::string> dir;

    dir['A'] = std::string("south");
    dir['B'] = std::string("north");
    dir['C'] = std::string("east");
    dir['D'] = std::string("west");

    // C++03 standard map iteration
    // This is more cumbersome, but shows better what is going on inside the loop.
    for (std::map<char, std::string>::iterator i = dir.begin(); i != dir.end(); i++)
        std::cout << "d[" << i->first << "] = " << i->second << std::endl;

    return 0;
}
```

Iterator is a pointer, so here we are using pointer dereferencing operator `->` to access `first` and `second` component of the map. Expression `i->first` is equivalent to `(*i).first` - that is dereferencing pointer `i`, and then accessing element `first` of the object pointed by `i`.

Output:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map3.cpp -o src/map3
$ ./src/map3
d[A] = south
d[B] = north
d[C] = east
d[D] = west
```

Keys not in the map

src/map4.cpp:

```
#include <iostream>
#include <map>

int main()
{
    std::map<char, std::string> dir;

    dir['A'] = std::string("north");
    dir['B'] = std::string("east");
    dir['C'] = std::string("south");
    dir['D'] = std::string("west");

    // Map size = 4
    std::cout << "dir.size() = " << dir.size() << std::endl;

    // Try to access value with key 'G' (creates new map entry with key 'G').
```

```

std::cout << "dir[G] = " << dir['G'] << std::endl;

// Map size = 5
std::cout << "dir.size() = " << dir.size() << std::endl;

return 0;
}

```

What happens when we look up the key that isn't present, is that a new key-value pair is stored in our container, where the value chosen is that of a "default". Since this map associates characters with strings, and `strings` have a default constructor, namely one that creates an empty string, we get that `dir['G']` takes on an empty string value. Output:

```

$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map4.cpp -o src/map4
$ ./src/map4
dir.size() = 4
dir[G] =
dir.size() = 5

```

Method `at()` and `map` container

If we want looking up an unknown key to result in an error, we can use `at`.

`src/map5.cpp`:

```

#include <iostream>
#include <map>

int main()
{
    std::map<char, std::string> dir;

    dir['A'] = std::string("north");
    dir['B'] = std::string("east");
    dir['C'] = std::string("south");
    dir['D'] = std::string("west");

    // Map size = 4
    std::cout << "dir.size() = " << dir.size() << std::endl;

    // Throws an exception -- out of range
    std::cout << "dir[G] = " << dir.at('G') << std::endl;

    return 0;
}

```

Output:

```

$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map5.cpp -o src/map5
$ ./src/map5
dir.size() = 4
dir.at(5) =
libc++abi.dylib: terminating with uncaught exception of type std::out_of_range:
map::at: key not found

```

Testing for a key

It's a bit funny, but we use the method `count` to test for existence of a key in a dictionary. Note that keys are unique, and so `count` can only return at most 1.

src/map6.cpp:

```
#include <iostream>
#include <map>

int main()
{
    std::map<char, std::string> dir;

    dir['A'] = std::string("north");
    dir['B'] = std::string("east");
    dir['C'] = std::string("south");
    dir['D'] = std::string("west");

    std::cout << "dir.count(A) = " << dir.count('A') << std::endl;
    std::cout << "dir.count(G) = " << dir.count('G') << std::endl;

    return 0;
}
```

Again, to emphasize: since keys are unique, the output of `count` can be either zero (key not found) or one (key found). Sample output of `src/map6.cpp` is:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map6.cpp -o src/map6
$ ./src/map6
dir.count(A) = 1
dir.count(G) = 0
```

Testing for a key

We can also use iterators! (They're quite handy, as you can see)

src/map7.cpp:

```
#include <iostream>
#include <map>

int main() {
    std::map<char, std::string> dir;

    dir['A'] = std::string("north");
    dir['B'] = std::string("east");
    dir['C'] = std::string("south");
    dir['D'] = std::string("west");

    char key = 'C';
    auto iter = dir.find(key); // Returns std::map<char, std::string>::iterator, but shorter to just use au
    if (iter == dir.end()) {
        std::cout << "key " << key << " is not present" << std::endl;
    }
    else {
        std::cout << "key " << key << " is present" << std::endl;
        std::cout << "value is " << iter->second << std::endl;
    }
}
```

```
    return 0;
}
```

Output:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map7.cpp -o src/map7
$ ./src/map7
key C is present
value is south
```

Key order

A map in C++ is by default an *ordered* collection (there is an unordered analogue). The consequence of this is that looking up an element costs logarithmic time with respect to the container size, as opposed to a constant time lookup afforded to us with an unordered container.

src/map8.cpp:

```
#include <iostream>
#include <map>

int main()
{
    std::map<char, std::string> dir;

    dir['C'] = std::string("south");
    dir['D'] = std::string("west");
    dir['B'] = std::string("east");
    dir['A'] = std::string("north");

    for (const auto& d : dir)
        std::cout << d.first << std::endl;

    return 0;
}
```

Keys of a map are always sorted. The order in which key-value pairs are added to the map does not matter.

Output:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map8.cpp -o src/map8
$ ./src/map8
A
B
C
D
```

Map and tuples

We can map strings to tuples. Going back to our census name data files:

src/map9.cpp:

```
#include <fstream>
#include <iostream>
#include <map>
#include <string>
#include <tuple>
```

```

int main() {
    // Open file and check if successful, print error message if it fails
    std::ifstream f("../dist.female.first");
    if (not f.good()) {
        std::cerr << "ERROR: Failed to open file" << std::endl;
        return 1;
    }

    // Create map 'names'
    std::map<std::string, std::tuple<double, double, int> > names;

    // Load file entries into the map
    std::string name;
    double perc1, perc2;
    int rank;
    while(f >> name >> perc1 >> perc2 >> rank) {
        names[name] = std::make_tuple(perc1, perc2, rank);
    }

    // Read from the map and print on std output
    // Method std::get<0>() gets 0th element of the tuple
    // The template parameter <0> must be a literal!
    // The reason is that we must be able to determine at compile time
    // which getter method to use.
    for(const auto& data : names) {
        std::cout << data.first << " " << std::get<2>(data.second) << std::endl;
    }

    return 0;
}

```

Here again we emphasize that when we use `std::get<i>`, the argument `i` must be a literal capable of being determined at compile time.

File `dist.female.first`:

MARY	2.629	2.629	1
PATRICIA	1.073	3.702	2
LINDA	1.035	4.736	3
BARBARA	0.980	5.716	4
ELIZABETH	0.937	6.653	5
JENNIFER	0.932	7.586	6
MARIA	0.828	8.414	7
TERRY	0.794	9.209	8
MARGARET	0.768	9.976	9
DOROTHY	0.727	10.703	10

Output:

```

$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/map9.cpp -o src/map9
$ ./src/map9
BARBARA 4
DOROTHY 10
ELIZABETH 5
JENNIFER 6
LINDA 3
MARGARET 9
MARIA 7
MARY 1

```


PATRICIA 2
TERRY 8

Using functions

We can decompose our functions into prototypes (header-files) and definitions (source-code files). Note the header guard that's used to protect us from re-declaring the same prototype.

src/readnames.hpp:

```
#ifndef READNAMES_HPP
#define READNAMES_HPP
```

```
#include <map>
#include <string>
#include <tuple>
```

```
std::map<std::string, std::tuple<double, double, int> > ReadNames(std::string filename);
```

```
#endif /* READNAMES_HPP */
```

src/readnames.cpp:

```
#include <fstream>
#include <iostream>
```

```
#include "readnames.hpp"
```

```
std::map<std::string, std::tuple<double, double, int>> ReadNames(std::string filename)
{
```

```
    // Create file I/O stream
```

```
    std::ifstream f(filename);
```

```
    // Create map 'names'
```

```
    std::map<std::string, std::tuple<double, double, int> > names;
```

```
    std::string name;
```

```
    double perc1, perc2;
```

```
    int rank;
```

```
    // Read file entries and store them into the map 'names'
```

```
    while(f >> name >> perc1 >> perc2 >> rank) {
```

```
        names[name] = std::make_tuple(perc1, perc2, rank); // Function that creates a tuple
```

```
    }
```

```
    // Return map 'filename' by value
```

```
    return names;
```

```
}
```

Alternative to a header guard, you could use `#pragma once`: which means to only include this file once (not standard)

src/testname.hpp:

```
#pragma once
```

```
#include <map>
#include <string>
#include <tuple>
```

```
double TestName(std::map<std::string, std::tuple<double, double, int>> names,
                std::string name);

src/testname.cpp:

#include <iostream>

#include "testname.hpp"

double TestName(std::map<std::string, std::tuple<double, double, int>> names,
                std::string name)
{
    // Variable to store name rank
    int name_rank = 0;

    // The variable 'match' is a map iterator. Function 'find(mapKey)' returns
    // the iterator that points to the map entry with key value 'mapKey'
    auto match = names.find(name);

    // Check if the iterator returns end value (i.e. 'mapKey' is not in the map).
    // If not, read the name rank for the 'name'.
    if (match != names.end())
    {
        // The name rank is the third entry (index 2) in the tuple 'match->second'.
        // It is retrieved by calling std::get<2> function.
        name_rank = std::get<2>(match->second);
    }

    return name_rank;
}
```

Using functions

```
src/main.cpp:

#include <iostream>
#include <string>
#include <vector>

#include "readnames.hpp"
#include "testname.hpp"

int main()
{
    // Read file and store its data in object 'names'.
    // Let compiler find the type of the object.
    auto names = ReadNames("../dist.female.first");

    // Create a vector of strings.
    std::vector<std::string> tests;
    tests.push_back("LINDA");
    tests.push_back("PETER");
    tests.push_back("DOROTHY");

    // Check for each name in the vector if it is stored in object 'names'.
    // If the name is found in object 'names' print its rank, otherwise print zero.
    for(auto test : tests)
```

```

{
    std::cout << test << " " << TestName(names, test) << std::endl;
}

return 0;
}

```

Output:

```

$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/main.cpp src/readnames.cpp \
$ src/testname.cpp -o src/main
$ ./src/main
LINDA 3
PETER 0
DOROTHY 10

```

Sets

This is also an *ordered* collection by default in C++ (which means logarithmic time lookups with respect to the input container size).

src/set.cpp:

```

#include <algorithm>
#include <fstream>
#include <iostream>
#include <set>
#include <string>

// Open file and copy its content into a set of strings
std::set<std::string> ReadNames(std::string filename)
{
    // Create a set of strings
    std::set<std::string> names;

    std::ifstream f(filename);
    if (not f.is_open())
    {
        std::cerr << "ERROR: Could not read file " << filename << std::endl;
        return names;
    }

    std::string name;
    double perc1, perc2;
    int rank;
    // Read file
    while (f >> name >> perc1 >> perc2 >> rank)
    {
        // Insert 'name' into the set, throw away other stuff
        names.insert(name);
    }
    f.close();

    // Return set of strings
    return names;
}

int main()

```

```

{
    // Create set of female names
    std::set<std::string> fnames = ReadNames("../dist.female.first");
    // Create set of male names
    std::set<std::string> mnames = ReadNames("../dist.male.first");

    // Create set of strings 'common' to store the intersection
    std::set<std::string> common; // Default set constructor

    // For more algorithms see http://en.cppreference.com/w/cpp/algorithm
    // Here we use set intersection algorithm...
    // std::inserter(c, i) function template is used to inserts an element
    // into container c at the iterator position i.
    // Returns std::insert_iterator
    // See: http://en.cppreference.com/w/cpp/iterator/inserter
    std::set_intersection(fnames.begin(),
                          fnames.end(),
                          mnames.begin(),
                          mnames.end(),
                          std::inserter(common, common.begin()));

    std::cout << fnames.size() << " female names" << std::endl;
    std::cout << mnames.size() << " male names" << std::endl;
    std::cout << common.size() << " common names" << std::endl;

    return 0;
}

```

The file `dist.female.names` is the same file used in the `map9.cpp` example. The file `dist.male.names` looks like this:

JUSTIN	0.311 49.040	56
TERRY	0.311 49.351	57
GERALD	0.309 49.660	58
KEITH	0.308 49.968	59
SAMUEL	0.306 50.274	60
WILLIE	0.302 50.576	61
RALPH	0.282 50.859	62
LAWRENCE	0.282 51.141	63
NICHOLAS	0.275 51.415	64
ROY	0.273 51.688	65

Output:

```

$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/set.cpp -o src/set
$ ./src/set
10 female names
10 male names
1 common names

```

Additional data structures

- `std::array` (C++ 2011)
- `std::list`
- `std::forward_list` (C++ 2011)

- `std::unordered_map` (C++ 2011)
- `std::unordered_set` (C++ 2011)

Note that `unordered_map` and `unordered_set` require the key to be hashable but in return we get constant time look-ups. On the other hand, the ordered analogues require a comparator operator to be defined for the key... more on this in 212!

Array example

Separate from a `vector`, there is also an `array` that is basically like a static array with methods. It's type and size is fixed at compile time.

src/array.cpp:

```
#include <array>
#include <iostream>

int main()
{
    std::array<double,4> a;

    a.fill(1.);
    a[2] = 3.;

    for (auto val : a)
        std::cout << val << std::endl;

    return 0;
}
```

Notice that we've used the `fill` method to initialize our container with ones. Output:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/array.cpp -o src/array
$ ./src/array
1
1
3
1
```


Linked lists

- Ordered data sequence similar to a C++ `vector` or Python list, but data is not stored contiguously.
- The access to individual list elements is maintained by links. Traversing to the *i*th element requires linear time work.
- There is additional storage overhead for the links.
- But this allows for insertion and removal operations in constant time, subject to having a pointer to the position/element we wish to insert into (or delete).


List example

src/list.cpp:

```
#include <iostream>
#include <list>
```



Singly linked list



Doubly linked list

Figure 1: fig

```
int main()
{
    // Create and populate list 'lst'.
    std::list<int> lst;
    lst.push_back(42);
    lst.push_back(17);
    lst.push_back(9);
    lst.push_front(18);

    // Print elements of the list.
    std::cout << "Elements of the list:\n";
    for (auto& val : lst)
        std::cout << val << std::endl;
    std::cout << "\n";

    // Create a list iterator and set it to the beginning of the list.
    auto it = lst.begin();

    // Advance list iterator to the third element of the list and erase it.
    // See: https://en.cppreference.com/w/cpp/iterator/advance
    // (remember 0-based indexing).
    advance(it, 2);
    std::cout << "Erasing element " << *it << " ... \n";
    // Dereference 'it' to get value ^^^
    lst.erase(it);
    std::cout << "\n";

    // Print elements of the list again to see the modified list.
    std::cout << "Elements of the list:\n";
    for (auto val : lst)
        std::cout << val << std::endl;
    std::cout << "\n";
}
```

```

    return 0;
}

$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/list.cpp -o src/list
$ ./src/list
Elements of the list:
18
42
17
9

Erasing element 17 ...

Elements of the list:
18
42
9

```

Recap: difference between (unordered) maps and sets

Maps and sets

- Python dictionaries and sets are internally implemented by using hashing.
- For hashing implementation, time complexity for data access is (amortized) constant time.
- Instances of C++ `std::map` and `std::set` are internally implemented using a tree data structure.
- For a tree, time complexity for data access is $O(\log n)$.
- Reference: <http://www.cplusplus.com/reference/map/map/operator%5B%5D/>.

Unordered maps and sets

- In the C++ 2011 standard the `std::unordered_map` and `std::unordered_set` were added.
- Like Python, internal implementation is based on hashing.
- Faster access, but entries are no longer ordered (that usually doesn't matter).

Unordered map example

```

src/unordered_map.cpp:

#include <iostream>
#include <unordered_map>

int main()
{
    std::unordered_map<int, std::string> dir;

    dir[0] = std::string("north");
    dir[1] = std::string("east");
    dir[2] = std::string("south");
    dir[3] = std::string("west");

    std::cout << "dir[2] = " << dir[2] << std::endl;
    std::cout << "dir[0] = " << dir[0] << std::endl;
}

```

```
    return 0;
}
```

Output:

```
$ clang++ -std=c++11 -Wall -Wextra -Wconversion src/unordered_map.cpp -o src/unordered_map
$ ./src/unordered_map
dir[2] = south
dir[0] = north
```

Reading

- **C++ Primer, Fifth Edition** by Lippman et al.
- Chapter 11: Associative Containers: Sections 11.1 - 11.3