

# Lecture 14

Fall 2020

## CME 211: Lecture 14

Topics:

- Compilation process
- Make for building software
- Debuggers

### Compilation

Although you can go from source code to an executable in one command, the process is actually made up of 4 steps

- Preprocessing
- Compilation
- Assembly
- Linking

`g++` (and `gcc` for C code) are driver programs that invoke the appropriate tools to perform these steps.

This is a high level overview. The compilation process also includes optimization phases during compilation and linking, and we'll have a lecture on this in CME212.

### Behind the scenes

We can inspect the compilation process in more detail with the `-v` compiler argument. `-v` typically stands for “verbose”.

Output:

```
$ g++ -v -Wall -Wextra -Wconversion src/hello1.cpp -o src/hello1
Apple LLVM version 7.3.0 (clang-703.0.31)
Target: x86_64-apple-darwin15.6.0
Thread model: posix
InstalledDir: /Applications/Xcode.app/Contents/Developer/Toolchains/XcodeDefault.xctoolchain/usr/bin
"/Applications/Xcode.app/Contents/Developer/Toolchains/XcodeDefault.xctoolchain/usr/bin/clang" -cc1 -trip
clang -cc1 version 7.3.0 (clang-703.0.31) default target x86_64-apple-darwin15.6.0
ignoring nonexistent directory "/usr/include/c++/v1"
#include "... " search starts here:
#include <...> search starts here:
/Applications/Xcode.app/Contents/Developer/Toolchains/XcodeDefault.xctoolchain/usr/bin/../include/c++/v1
/usr/local/include
/Applications/Xcode.app/Contents/Developer/Toolchains/XcodeDefault.xctoolchain/usr/bin/../lib/clang/7.3.0
/Applications/Xcode.app/Contents/Developer/Toolchains/XcodeDefault.xctoolchain/usr/include
/usr/include
```

```

/System/Library/Frameworks (framework directory)
/Library/Frameworks (framework directory)
End of search list.
"/Applications/Xcode.app/Contents/Developer/Toolchains/XcodeDefault.xctoolchain/usr/bin/ld" -demangle -dy

```

## Splitting up the steps manually

GNU compiler flags:

- -E: preprocess
- -S: compile
- -c: assemble

Output:

```

$ cat src/hello1.cpp
#include <iostream>

int main() {
    std::cout << "Hello, CME 211!" << std::endl;
    return 0;
}
$ g++ -E -o src/hello1.i src/hello1.cpp
$ g++ -S -o src/hello1.s src/hello1.i
clang: warning: treating 'cpp-output' input as 'c++-cpp-output' when in C++ mode, this behavior is deprecated
$ g++ -c -o src/hello1.o src/hello1.s
$ g++ -o src/hello1 src/hello1.o
$ ./src/hello1
Hello, CME 211!

```

## Preprocessing

The preprocessor handles the lines that start with #:

- #include
- #define
- #if
- etc.

You can invoke the preprocessor with the `cpp` command.

## Preprocessed file

From `src/hello1.i`:

```

# 1 "hello1.cpp"
# 1 "<command-line>"
# 1 "/usr/include/stdc-predef.h" 1 3 4
# 31 "/usr/include/stdc-predef.h" 2 3 4
// ... a bunch of omitted lines
namespace std {
    // We'll learn what the following lines mean in 212.
    typedef long unsigned int size_t;
    typedef long int ptrdiff_t;

    typedef decltype(nullptr) nullptr_t;
}

```

```
// approximately 17,500 more lines omitted!
```

```
int main() {  
    std::cout << "Hello" << std::endl;  
    return 0;  
}
```

If you're curious about what the first few lines beginning with `#` signs represent, see the documentation: <https://gcc.gnu.org/onlinedocs/gcc-4.8.5/cpp/Preprocessor-Output.html>. "Source file name and line number information is conveyed by lines of the form

```
# linenum filename flags
```

These are called *linemarkers*... They mean that the following line originated in file *filename* at line *linenum*... After the file name comes zero or more flags, which are '1', '2', '3', or '4'. If there are multiple flags, spaces separate them. Here is what the flags mean..."

## Compilation

- Compilation is the process of translating source code (i.e. the C++ code you wrote) into assembly.
- The assembly commands are still human readable text (if the human knows assembly)!

Note that we could look at `src/hello.s`, but because we are using a library `iostream` the assembly commands become a bit harder to interpret (you can look at them on your own if you wish). Instead we'll turn to a simple addition function file: `src/add.cpp`.

```
#include <iostream>  
  
int add(int a, int b) {  
    return a + b;  
}  
  
int main(int argc, char* argv[]) {  
    int a, b;  
    a = atoi(argv[1]);  
    b = atoi(argv[2]);  
    int c = add(a, b);  
    std::cout << c << std::endl;  
    return 0;  
}
```

We can run compilation up through assembly by invoking `g++ -S -o src/add.s src/add.cpp`, and we can inspect a few key snippets. Let's first look at the addition procedure, i.e. our `add` function:

```
_Z3addii:  
.LFB1493:  
    .cfi_startproc  
    pushq   %rbp  
    .cfi_def_cfa_offset 16  
    .cfi_offset 6, -16  
    movq    %rsp, %rbp  
    .cfi_def_cfa_register 6  
    movl    %edi, -4(%rbp)  
    movl    %esi, -8(%rbp)  
    movl    -4(%rbp), %edx  
    movl    -8(%rbp), %eax  
    addl    %edx, %eax  
    popq    %rbp
```

Next let's see how our function gets invoked; we'll skip most of the output but print a few key lines:

```
main:
.LFB1494:
    .cfi_startproc
    pushq   %rbp
    ...
    movq    -32(%rbp), %rax    <-- Here we read the first argument from command line.
    addq    $8, %rax          <-- We have an offset from our char* array.
    movq    (%rax), %rax
    movq    %rax, %rdi
    call    atoi@PLT
    movl    %eax, -12(%rbp)
    movq    -32(%rbp), %rax    <-- Here we read the second argument from command line.
    addq    $16, %rax          <-- Note the different offset.
    movq    (%rax), %rax
    movq    %rax, %rdi
    call    atoi@PLT
    movl    %eax, -8(%rbp)     <-- Here we set up our call to add.
    movl    -8(%rbp), %edx
    movl    -12(%rbp), %eax
    movl    %edx, %esi
    movl    %eax, %edi
    call    _Z3addii           <-- Invoke add operator.
    movl    %eax, -4(%rbp)
    movl    -4(%rbp), %eax
    movl    %eax, %esi
    ...
    movl    $0, %eax          <-- Return 0 from main function.
```

## Assembly

This step translates the human-readable assembly into binary machine code in a `.o` file.

- `.o` files are called object files
- Linux uses the Executable and Linkable Format (ELF) for these files
- If you try to look at these files with a normal text editor you will just see garbage, intermixed with a few strings
- Sometimes it is helpful to inspect object files with the `nm` command to see what symbols are defined:

Output:

```
$ nm ./src/hello1.o
... <-- Output omitted
__ZNSt3__112basic_stringIcNS_11char_traitsIcEENS_9allocatorIcEEE6__initEmc
        U __ZNSt3__112basic_stringIcNS_11char_traitsIcEENS_9allocatorIcEEED1Ev
        U __ZNSt3__113basic_ostreamIcNS_11char_traitsIcEEE3putEc
        U __ZNSt3__113basic_ostreamIcNS_11char_traitsIcEEE5flushEv
... <-- Output omitted
0000000000000000 T _main
        U _memset
        U _strlen
```

Here is some documentation. The notable aspects are that there is a symbol associated with each subroutine. E.g. the `T _main` indicates that `main()` is a global text symbol, whereas `U _memset` indicates that this function call currently contains an unresolved reference (to be resolved by the linker in the subsequent step).

## Linking

- Linking is the process of building the final executable by combining (linking) the `.o` file(s), and possibly library files as well
- The linker makes sure all of the required functions are present
- If for example `foo.o` contains a call to a function called `bar()`, there has to be another `.o` file or library file that provides the implementation of the `bar()` function

### Linking example

src/foobar.hpp:

```
#pragma once
```

```
void bar(void);  
void foo(void);
```

src/foo.cpp:

```
#include <iostream>
```

```
void foo(void) {  
    std::cout << "Hello from foo" << std::endl;  
}
```

src/bar.cpp:

```
#include <iostream>
```

```
void bar(void) {  
    std::cout << "Hello from bar" << std::endl;  
}
```

src/main.cpp:

```
#include "foobar.hpp"
```

```
int main() {  
    foo();  
    bar();  
    return 0;  
}
```

### Linking example

Inspect the files:

Output:

```
$ ls src/foo* src/bar* src/foobar* src/main*  
src/bar.cpp src/foobar.hpp src/foobar.hpp src/foo.cpp src/main.cpp
```

Compile and assemble source files, but don't link: Output:

```
$ g++ -c src/foo.cpp -o src/foo.o  
$ g++ -c src/bar.cpp -o src/bar.o  
$ g++ -c src/main.cpp -o src/main.o
```

Let's inspect the output: Output:

```
$ ls src/*.o
src/bar.o  src/foo.o  src/main.o
```

What symbols are present in the object files?

Output:

```
$ nm src/foo.o
... <-- Output omitted
      U __ZNSt3__112basic_stringIcNS_11char_traitsIcEENS_9allocatorIcEEE6__initEmc
      U __ZNSt3__112basic_stringIcNS_11char_traitsIcEENS_9allocatorIcEEED1Ev
      U __ZNSt3__113basic_ostreamIcNS_11char_traitsIcEEE3putEc
      U __ZNSt3__113basic_ostreamIcNS_11char_traitsIcEEE5flushEv
      U __ZNSt3__113basic_ostreamIcNS_11char_traitsIcEEE6sentryC1ERS3_
      U __ZNSt3__113basic_ostreamIcNS_11char_traitsIcEEE6sentryD1Ev
00000000000000630 S __ZNSt3__116__pad_and_outputIcNS_11char_traitsIcEEEENS_19ostreambuf_iteratorIT_TO_EES6_
00000000000001a0 S __ZNSt3__124__put_character_sequenceIcNS_11char_traitsIcEEEERNS_13basic_ostreamIT_TO_EE
      U __ZNSt3__14coutE
0000000000000090 S __ZNSt3__14endlIcNS_11char_traitsIcEEEERNS_13basic_ostreamIT_TO_EES7_
... <-- Output omitted.
      U _memset
      U _strlen

$ nm src/bar.o
... <-- Output omitted.
      U __ZNSt3__113basic_ostreamIcNS_11char_traitsIcEEE3putEc
      U __ZNSt3__113basic_ostreamIcNS_11char_traitsIcEEE5flushEv
... <-- Output omitted.
      U _memset
      U _strlen

$ nm src/main.o
      U __Z3barv
      U __Z3foov
0000000000000000 T _main
```

Notice that in the last step we still have undefined references to `foo` and `bar`. What happens if we try to link `main.o` into an executable with out pointing to the other object files?

Output:

```
$ g++ src/main.o -o src/main
Undefined symbols for architecture x86_64:
  "bar()", referenced from:
      _main in main.o
  "foo()", referenced from:
      _main in main.o
ld: symbol(s) not found for architecture x86_64
clang: error: linker command failed with exit code 1 (use -v to see invocation)
```

Ahhh, linker errors! Let's do it right, by giving all of the information that the main program needs in order to execute on the instructions. Output:

```
$ g++ src/main.o src/foo.o src/bar.o -o src/main
$ ./src/main
Hello from foo
Hello from bar
```

## Libraries

- Libraries are really just a file that contain one or more `.o` files

- On Linux these files typically have a `.a` (static library) or `.so` (dynamic library) extension
- `.so` files are analogous to `.dll` files on Windows
- `.dylib` files on Mac OS X and iOS are also very similar to `.so` files
- Static libraries are factored into the executable at link time in the compilation process.
- Shared (dynamic) libraries are loaded up at run time.

## JPEG Example

From `src/hw6.cpp`:

```
// code omitted

#include <jpeglib.h>

#include "hw6.hpp"

void ReadGrayscaleJPEG(std::string filename, boost::multi_array<unsigned char,2> &img)
{
    /* Open the file, read the header, and allocate memory */

    FILE *f = fopen(filename.c_str(), "rb");
    if (not f)
    {
        std::stringstream s;
        s << __func__ << ": Failed to open file " << filename;
        throw std::runtime_error(s.str());
    }
    // code omitted
}

// code omitted

#ifdef DEBUG
int main()
{
    boost::multi_array<unsigned char,2> img;
    ReadGrayscaleJPEG("stanford.jpg", img);
    WriteGrayscaleJPEG("test.jpg", img);

    return 0;
}
#endif /* DEBUG */
```

Let's try to compile:

Output:

```
$ g++ -std=c++11 -Wall -Wextra -Wconversion src/hw6.cpp -o src/hw6
```

```
Undefined symbols for architecture x86_64:
```

```
  "_jpeg_CreateCompress", referenced from:
```

```
      WriteGrayscaleJPEG(std::__1::basic_string<char, std::__1::char_traits<char>, std::__1::allocator<char>>, boost::multi_array<unsigned char, 2>, unsigned int, unsigned int)
```

```
...  <-- Output omitted.
```

```
  "_jpeg_write_scanlines", referenced from:
```

```
      WriteGrayscaleJPEG(std::__1::basic_string<char, std::__1::char_traits<char>, std::__1::allocator<char>>, boost::multi_array<unsigned char, 2>, unsigned int, unsigned int)
```

```
  "_main", referenced from:
```

```
implicit entry/start for main executable
ld: symbol(s) not found for architecture x86_64
clang: error: linker command failed with exit code 1 (use -v to see invocation)
```

That did not work. The linker looks for the `main` symbol when trying to build and executable. This linker also cannot find all of the symbols from the JPEG library.

Let's find the `jpeglib.h` header file:

Output:

```
$ locate jpeglib.h
/usr/local/Cellar/jpeg/8d/include/jpeglib.h
/usr/local/include/jpeglib.h      <-- We're going to link to this file.
```

Let's find `libjpeg`, it's the library that actually contains the `jpeglib.h` header file.

```
$ locate libjpeg
/Applications/Xcode.app/Contents/Applications/Application Loader.app/Contents/itms/java/lib/libjpeg.dylib
/usr/local/Cellar/jpeg/8d/lib/libjpeg.8.dylib
/usr/local/Cellar/jpeg/8d/lib/libjpeg.a
/usr/local/Cellar/jpeg/8d/lib/libjpeg.dylib
/usr/local/Homebrew/Library/Taps/homebrew/homebrew-core/Aliases/libjpeg
/usr/local/Homebrew/Library/Taps/homebrew/homebrew-core/Aliases/libjpeg-turbo
/usr/local/lib/libjpeg.8.dylib
/usr/local/lib/libjpeg.a          <-- We're going to link to this file.
/usr/local/lib/libjpeg.dylib
/usr/local/lib/python3.5/site-packages/PIL/.dylibs/libjpeg.9.dylib
```

Note that the library files may be in a different location on your system.

Now let's compile:

Output:

```
$ g++ -std=c++11 -Wall -Wextra -Wconversion src/hw6.cpp -o src/hw6 -DDEBUG \
-I/usr/local/include -L/usr/local/lib -ljpeg
$ ./src/hw6
```

- `-I/usr/local/include`: look in this directory for include files (optional in this case)
- `-L/usr/local/lib`: look in this directory for library files (optional in this case, maybe required on Ubuntu)
- `-ljpeg`: link to the `libjpeg.{a,so}` file (not optional here)

## Make

- Utility that compiles programs based on rules read in from a file called Makefile
- Widely used on Linux/Unix platforms
- Setup and maintenance of Makefile(s) can become rather complicated for major projects
- We will look at a few simple examples

### Example source files

`src/ex1/sum.cpp`:

```
#include "sum.hpp"

double sum(double a, double b) {
    double c = a + b;
```



```

    return c;
}

src/ex1/sum.hpp:

#pragma once

double sum(double a, double b);

src/ex1/main.cpp:

#include <iostream>

#include "sum.hpp"

int main() {
    double a = 2., b = 3., c;

    c = sum(a,b);
    std::cout << "c = " << c << std::endl;

    return 0;
}

```

### Example makefile

```

src/ex1/makefile:

main: main.cpp sum.cpp sum.hpp
    g++ -Wall -Wextra -Wconversion -o main main.cpp sum.cpp

Anatomy of a make rule:

target: dependencies
    build_command

```

- **target:** is the thing you want the rule to create. The target should be a file that will be created in the file system. For example, the final executable or intermediate object file.
- **dependencies:** space separated list files that the target depends on (typically source or header files)
- **build\_command:** a **tab-indented** shell command (or sequence) to build the target from dependencies.

### Let's run the example

Let's run make! From `src/ex1/`:

```

$ ls
main.cpp  makefile  sum.cpp  sum.hpp
$ make
g++ -Wall -Wextra -Wconversion -o main main.cpp sum.cpp
$ ls
main  main.cpp  makefile  sum.cpp  sum.hpp
$ make
make: 'main' is up to date.
$

```

### File changes

Make looks at time stamps on files to know when changes have been made and will recompile accordingly:

```
$ make
make: 'main' is up to date.
$ touch main.cpp
$ make
g++ -Wall -Wextra -Wconversion -o main main.cpp sum.cpp
$ touch sum.hpp
$ make
g++ -Wall -Wextra -Wconversion -o main main.cpp sum.cpp
$ make
make: 'main' is up to date.
```

## Make variables, multiple targets, and comments

src/ex2/makefile:

```
# this is a makefile variable, note := for direct assignment
CXX := g++

# This line and the next two are makefile comments.
#CXXFLAGS := -Wall -Wextra -Wconversion
#CXXFLAGS := -Wall -Wextra -Wconversion -g
CXXFLAGS := -Wall -Wextra -Wconversion -fsanitize=address

# We use $(VAR) to enable parameter expansion.
main: main.cpp sum.cpp sum.hpp
    $(CXX) $(CXXFLAGS) -o main main.cpp sum.cpp

# here is a target to clean up the output of the build process
.PHONY: clean
clean:
    $(RM) main
```

The last target is a PHONY one because it doesn't produce any files. Output (from `src/ex2` directory):

```
$ ls
main.cpp makefile sum.cpp sum.hpp
$ make
g++ -Wall -Wextra -Wconversion -fsanitize=address -o main main.cpp sum.cpp
$ ls
main main.cpp makefile sum.cpp sum.hpp
$ make clean
rm -f main
$ ls
main.cpp makefile sum.cpp sum.hpp
```

## Individual compilation of object files

Make has automatic variables such as `$$` and `$$<`, where the former specifies the name of the target of the rule, and the latter specifies the name of the first pre-requisite.

src/ex3/makefile:

```
CXX := g++
CXXFLAGS := -O3 -Wall -Wextra -Wconversion -std=c++11

TARGET := main
OBJS := main.o sum.o foo.o bar.o
INCS := sum.hpp foobar.hpp
```

```
$(TARGET): $(OBJS)
    $(CXX) -o $(TARGET) $(OBJS)

# this is a make pattern rule
%.o: %.cpp $(INCS)
    $(CXX) -c -o $@ $< $(CXXFLAGS)

.PHONY: clean
clean:
    $(RM) $(OBJS) $(TARGET)
```

Output (from src/ex3 directory):

```
$ ls
bar.cpp  foo.cpp  main.cpp  makefile  sum.cpp  sum.hpp
$ make
g++ -c -o main.o main.cpp -O3 -Wall -Wextra -Wconversion -std=c++11
g++ -c -o sum.o sum.cpp -O3 -Wall -Wextra -Wconversion -std=c++11
g++ -c -o foo.o foo.cpp -O3 -Wall -Wextra -Wconversion -std=c++11
g++ -c -o bar.o bar.cpp -O3 -Wall -Wextra -Wconversion -std=c++11
g++ -o main main.o sum.o foo.o bar.o
$ ls
bar.cpp  bar.o  foo.cpp  foo.o  main  main.cpp  main.o  makefile  sum.cpp  sum.hpp  sum.o
$ make clean
rm -f main.o sum.o foo.o bar.o main
$ ls
bar.cpp  foo.cpp  main.cpp  makefile  sum.cpp  sum.hpp
```

## Linking to a library & run targets

src/ex4/makefile:

```
# conventional variable for c++ compiler
CXX := g++

# conventional variable for C preprocessor
CPPFLAGS := -DDEBUG

# conventional variable for C++ compiler flags
CXXFLAGS := -O3 -std=c++11 -Wall -Wextra -Wconversion

# conventional variable for linker flags
LDFLAGS := -ljpeg

TARGET := hw6
OBJS := hw6.o
INCS := hw6.hpp

$(TARGET): $(OBJS)
    $(CXX) -o $(TARGET) $(OBJS) $(LDFLAGS)

%.o: %.cpp $(INCS)
    $(CXX) -c -o $@ $< $(CPPFLAGS) $(CXXFLAGS)

# use .PHONY for targets that do not produce a file
.PHONY: clean
```

```
clean:
    rm -f $(OBJS) $(TARGET) *~

.PHONY: run
run: $(TARGET)
    ./${TARGET}
```

Output (from src/ex4 directory):

```
$ ls
hw6.cpp hw6.hpp makefile stanford.jpg
$ make
g++ -c -o hw6.o hw6.cpp -DDEBUG -O3 -std=c++11 -Wall -Wextra -Wconversion
g++ -o hw6 hw6.o -ljpeg
$ ./hw6
$ make clean
rm -f hw6.o hw6 *~
$ make run
g++ -c -o hw6.o hw6.cpp -DDEBUG -O3 -std=c++11 -Wall -Wextra -Wconversion
g++ -o hw6 hw6.o -ljpeg
./hw6
$ ls
hw6 hw6.cpp hw6.hpp hw6.o makefile stanford.jpg test.jpg
```

## Make

- Automation tool for expressing how your C/C++/Fortran/LaTeX code should be built.
- Good for single platform projects.
- But be careful with dependencies. It is **very** important to understand this process for larger projects.
- Hand writing Makefile(s) for cross-platform projects is not recommended. You should consider using configuration tools such as CMake.